

Local Open-Weight LLMs Are Competitive with Commercial APIs for Political Science Text Classification

Hanno Hilbig*

2026-05-18

Abstract

Local open-weight models for LLM text classification have several advantages: they avoid API costs, keep data on the researcher’s machine, and make exact model versions easier to preserve. I benchmark five local models against four commercial API models on 34 political science classification tasks. Local models are competitive, especially on simpler tasks. Comparing the best local and best API model within each task, local models match or exceed API performance on 9 tasks. On average, the best API model exceeds the best local model by 0.015 on the main F1 score. The four strongest models, Claude Sonnet 4.6, gpt-5.5, Gemma 4 31B, and DeepSeek V4 Pro, fall within 0.021 mean F1. API models have their clearest edge on complex tasks with many labels or multiple outputs per item. Batching several items in one prompt usually reduces local runtime per item, but this sometimes returns invalid response formats. Overall, local open-weight models are a practical alternative for many political science classification tasks, provided researchers validate candidate models on task-specific labels and check batching reliability before scaling up.

1 Motivation

When using LLMs to classify text data, researchers can use commercial API models or local open-weight models. Commercial API models are typically reliable and easy to call from R or Python. However, they cost money per API call, send data off-machine, and create reproducibility risk when model behavior changes over time. The alternative is local open-weight models. These models are free to run, keep the data on the researcher’s machine, and make exact model versions easier to preserve and track. This matters when researchers work with restricted data or when data-use rules prohibit third-party APIs. Open-weight models can also be useful when research funds are not abundant, or for exploratory tasks that may otherwise not be pursued because of costs.

This benchmark focuses on three questions: whether local models approach commercial API performance, how long local classification takes, and when task complexity widens the API advantage. I use 34 tasks from published political science replication archives because they cover the kinds of coding problems applied researchers routinely face: relevance, stance, tone, events, claims, relations, issues, and topics.

The main result is practical rather than a single leaderboard. Local models are often close to API models, especially on simpler tasks, but performance varies enough across tasks that researchers should validate candidate models on the target coding problem. APIs retain their clearest edge on more complex tasks. Prompt batching can reduce local runtime, but only after checking invalid-output rates.

2 Design

The benchmark has two parts. First, I run a one-item-at-a-time benchmark covering thirty-four classification tasks drawn from published political science replication archives and nine models: five local open-weight

*hhilbig@ucdavis.edu. Code and data: <https://github.com/hhilbig/polsci-open-bench>.

models and four commercial API models from OpenAI, Anthropic, and DeepSeek. This produces 306 model-task comparisons. Each model sees the same items within a task. Thirty-one tasks use 500 sampled items. Three tasks have fewer cleaned items available, so I use all available items: 320 for Brandt political relevance, 312 for PLOVER CAMEO event coding, and 293 for COVID threat minimization. This is why task sizes range from 293 to 500 items.

I report performance and reliability for each task and model. The main F1 score puts precision and recall on the same scale. For binary tasks, I use the F1 score for the positive class, the harmonic mean of precision and recall. For multi-binary tasks, I average the per-label F1 scores. For single-label categorical tasks, I use macro F1, which gives each observed class equal weight before averaging. Accuracy is the share of items classified correctly. Matthews correlation coefficient (MCC) summarizes overall agreement while accounting for false positives and false negatives. I report all three because class imbalance is common: macro-style F1 highlights rare-class performance, accuracy gives the plain success rate, and MCC is less sensitive than accuracy to majority-class dominance.

I also record time per item and the share of unusable responses. A response is unusable if it fails to parse, returns the wrong object shape, or uses a label outside the allowed label set. I compute performance metrics over usable outputs only, so failure rates need to be read alongside F1, accuracy, and MCC.

I ran each model on each item once in the one-item-at-a-time benchmark. The local setup was a single Apple Silicon MacBook with 32 GB of unified memory running Ollama at temperature 0.1. Commercial API calls used provider-specific constrained-output settings. OpenAI tiers used `reasoning_effort=medium`; DeepSeek ran with internal thinking disabled; OpenAI and DeepSeek returned JSON-constrained outputs; and Anthropic used tool use to constrain labels. I use the source paper’s coding instructions verbatim where possible; otherwise, I derive prompts from the published codebook (see Table 1).

Second, I test whether prompt batching speeds up local inference. In the batching experiment, the model receives ten items in one prompt and returns ten labels in a single response. This tests whether batching reduces time per item without changing the classification task. The local grid with 10 items per prompt covers all thirty-four tasks and all five local models.

The one-item-at-a-time benchmark is the basis for all main performance comparisons. I treat prompt batching as a separate speed and reliability check because it changes the format of the model call and can change invalid-output rates. Separating the two designs keeps the performance comparison clean while still showing whether local use is practical at larger scale.

Finally, I use uncertainty checks to avoid overinterpreting small differences. Within-task comparisons use 95% paired-by-item bootstrap confidence intervals (1000 iterations). I treat the thirty-four tasks as a fixed benchmark population, but also report paired task-level bootstrap intervals that resample tasks with replacement.

3 Results

3.1 Average performance

Figure 1 shows the 34-task model averages and the spread of task-level performance within each model. Claude Sonnet 4.6 has the highest observed mean F1 (0.668), followed by gpt-5.5 (0.660), Gemma 4 31B (0.648), and DeepSeek V4 Pro (0.647). The spread across these four models is 0.021 points. The same broad top group appears under mean accuracy and mean MCC (Table 6), although the exact order changes slightly. Among the OpenAI tiers, gpt-5.4-nano remains the cost-efficient option, at 0.632 mean F1 compared with 0.660 for gpt-5.5.

The task-level bootstrap gives the same result (Table 3). The Claude Sonnet 4.6 versus gpt-5.5 interval includes zero, Gemma 4 31B and DeepSeek V4 Pro are effectively tied, and the estimated Claude advantage over Gemma 4 31B and DeepSeek V4 Pro is about two F1 points. Many within-task leader differences are also too uncertain to treat as firm rankings. In practice, these small average gaps are not enough to rank the models reliably.




Figure 1: Mean F1 across thirty-four tasks per model. Large filled circles show 34-task means. Small gray dots show observed per-task mean F1 values.

The local/API comparison in Figure 2 directly addresses the paper’s main question. After selecting the best model within each class for each task, the best local model matches or exceeds the best API model on 9 of the 34 tasks. On average, the best API model exceeds the best local model by 0.015 F1, and the median API advantage is 0.019. Figure 2 uses the same sign convention: values below zero favor local models and values above zero favor API models. This is not evidence that local models dominate API models. It shows that the API advantage is often small enough that local model calls are worth validating on the target task.

Table 4 reports how much performance would improve if researchers could choose the best model separately for each task. These upper bounds choose the best model after observing performance on each task. Allowing the local model to vary by task raises mean F1 from 0.648 for the best single local model to 0.673. Across all nine models, task-specific selection reaches 0.696, compared with 0.668 for the best single model. These values show why task-specific validation can be valuable, but they do not show how large a validation sample is needed to choose models reliably.

The larger differences are across tasks. The paired-by-item intervals are wide enough that many task-level leader differences remain uncertain. I therefore do not read the benchmark primarily as a leaderboard: global rank is a weak decision rule for applied coding.

3.2 Different models lead on different tasks

Figure 3 groups the thirty-four tasks into five broad annotation types: relevance and harm, position and tone, events and actions, claims and relations, and issues and topics. These are my groupings, not source-provided task families, and I assign borderline cases by the prediction target. The full task list is in Table 1. The type-level averages do not show a simple local versus API pattern.

At the task level, the top point estimates are split across eight models (Table 2). gpt-5.5 leads on thirteen tasks, Claude Sonnet 4.6 leads on eight, Mistral Small leads on three, and five other models lead on two tasks each.

The winner split matters because model strengths differ by task. Even models with lower average performance sometimes produce the best task-level point estimate. A model that is weaker on average can still be the best model for a specific task.




Figure 2: API advantage in main F1 by task. Values below zero favor the local model class; values above zero favor the API model class.

3.3 Performance is lower on more complex coding tasks

Figure 4 examines whether coding complexity explains the large task-to-task variation. I measure complexity with two features that researchers can observe before running a benchmark: the prompt/codebook and the number of labels used in the gold data. Effective labels measure how many labels are actually used in the gold data, rather than how many labels appear in the codebook. Low-complexity tasks have fewer than three effective labels and short prompts. Medium-complexity tasks have either more labels or longer prompts. High-complexity tasks have at least eight effective labels or require multi-binary outputs.

Performance falls as coding tasks become more complex, but local models remain close on many tasks. Mean F1 falls from 0.70 to 0.57 for API model-task pairs and from 0.67 to 0.52 for local pairs as tasks move from low to high complexity. When I instead compare the best API model with the best local model on each task, the low-complexity gap nearly disappears (0.740 versus 0.735), while the high-complexity gap is about 0.05 F1 (0.605 versus 0.555). This is the clearest API edge in the complexity groups: long codebooks, many effective labels, and tasks with multiple outputs per item.

Figure 5 shows one reason for the complexity result: tasks with more actively used labels tend to show larger API advantages. Effective label count measures how many labels actually matter in the sample, rather than simply counting every possible label. The relationship is positive. The simple correlation between effective label count and the API-minus-local gap is about 0.50. Binary and two-class tasks have a mean API advantage of 0.004, while many-class or multi-label tasks have a mean API advantage of 0.025.

3.4 Local inference is not uniformly slower

Local inference speed varies sharply with model size, model design, and batching (Figure 6). On a typical 500-item task, median one-item-at-a-time runtimes range from about 5 minutes for Qwen3-30B-A3B and 7 minutes for Gemma 4 26B to about 29 minutes for Gemma 4 31B. Qwen3-30B-A3B is fastest because it is




Figure 3: Mean F1 within five broad annotation types, by model. Type-level averages are descriptive only.

a sparse model that uses about three billion active parameters per call. Gemma 4 31B has the strongest local 34-task one-item-at-a-time F1, while Gemma 4 26B is much faster and remains close on average. Local inference is slow enough to plan around, but not slow enough to rule out ordinary validation and production coding.

All five local models ran on a single 32 GB Apple Silicon MacBook; the largest, Gemma 4 31B, fit in unified memory. Appendix Figure 7 translates the local runtimes into median minutes per 1,000 items for each model. For reference, the API tiers in this run took roughly 0.9 to 1.7 seconds per item end-to-end, but those numbers mix model speed with provider load, network delays, and queueing and are not shown in the figure.

3.5 Batching speeds up local inference but can produce invalid output

Prompt batching changes the prompt: the model receives several items at once and must return one label for each item. Provider Batch API mode is different. It changes how requests are queued and billed, but each item still appears in a separate request. I test prompt batching here. Pipal et al. (2026) show why this matters for labeling tasks: one-at-a-time classification repeats the same prompt and request setup for every item.

Figure 6 shows the main local result. Prompt batching reduces median runtime per item for every local model, with model-level median speedups from about $1.23\times$ to $1.86\times$. Average F1 changes are modest, ranging from -0.023 to $+0.010$. Gemma 4 26B is the most useful fast batched option in these results: it remains close to Gemma 4 31B on average but classifies a typical 500-item task in about 4 minutes with 10 items per prompt,




Figure 4: Coding complexity and model performance across the 34-task benchmark. The left panel shows mean F1 across all model-task results, separately for local and API models. The right panel shows the best local model and best API model per task within each complexity group. Points show group means; lines connect complexity groups within model class.




Figure 5: API advantage in main F1 by effective number of labels. The dashed horizontal line marks equal performance; the solid line is a descriptive linear fit across tasks.

compared with about 21 minutes for Gemma 4 31B.

Batching helps less when prompts are long. Short-prompt model-task pairs have a median local speedup of about $1.34\times$ with 10 items per prompt, while long-codebook pairs have a median speedup of about $1.20\times$. The bigger problem is invalid output: some models return invalid JSON, the wrong object shape, or labels outside the allowed set. The largest failures are concentrated in specific model-task pairs. For example, Gemma 4 26B produces unusable responses for 72% of Erlich ATI items with 10 items per prompt, and Qwen3-30B-A3B produces unusable responses for 51% of Halterman CCC items with 10 items per prompt. In a separate 10-task prompt-batching diagnostic for OpenAI models, gpt-5.5 returned unusable output for 68%




Figure 6: Mean F1 vs median runtime per item, for the five local models. Circles show one-at-a-time calls over the 34-task benchmark; triangles show calls with 10 items per prompt over the same task set. Lines connect the two call modes for the same model, and labels identify model pairs. API models are excluded because their per-call time mixes server load, queueing, and network delay and is not directly comparable to local inference time.

of Halterman CCC items. That OpenAI diagnostic does not enter the main one-item-at-a-time performance results. Table 5 lists the covered response-format and label failure rates above 5%.

3.6 Cost

The OpenAI portion of the 34-task one-item-at-a-time benchmark remains cheap enough for benchmarking on a validation sample, but the tiers differ sharply. At publicly listed prices, gpt-5.4-nano costs about \$4 for the 16,425 predictions in the 34-task benchmark, while gpt-5.5 would cost about \$71 under direct per-request pricing. In return, gpt-5.5 improves mean F1 from 0.632 to 0.660. The smaller tier is therefore worth testing before moving to a flagship API model.

Provider Batch API mode changes the cost calculation for flagship models. This is provider-side request processing, not prompt batching with several items in one prompt. For a 16-task subset, I ran gpt-5.5 through the OpenAI Batch API: 7,605 requests completed with 0 parse/API errors at an estimated batch-discounted cost of \$14.13 under a \$20 budget cap. I also ran Claude Sonnet 4.6 through provider Batch API mode for the same subset; 7,604 of 7,605 outputs parsed cleanly after retries. Local models avoid API charges, but the relevant cost becomes runtime, hardware availability, and energy use.

Implications for applied use

The results suggest a simple validation step rather than a single best model. Start with 100 to 300 hand-labeled cases from the target task. Test one cheap API model, one flagship API model, and at least two local models. Report the main F1 score, accuracy, MCC, and the share of unusable responses. For imbalanced tasks, inspect rare-class recall as well.

Researchers should test batching on the same validation sample before production coding. In these results, 10 items per prompt usually reduces local runtime with little average performance loss, but the failure rate can be high for specific model-task pairs. For long, multi-class codebooks, researchers should not batch aggressively unless retries are built in.

All five local models fit on a 32 GB Apple Silicon MacBook, so one laptop setup is enough hardware for the full benchmark. When paid API use is acceptable, cheap tiers should also be tested rather than dismissed: in

this run, gpt-5.4-nano lost 0.028 mean F1 relative to gpt-5.5 and cost about \$4 for the full 34-task benchmark.

Limitations and scope

Model selection and prompting. The model set reflects five specific local models that fit a 32 GB Apple Silicon setup, plus four commercial API tiers that include cheaper and flagship options. I did not optimize prompts separately per model. Some prompts use the source paper’s wording; others come from published codebooks.

Beyond this benchmark. Local inference keeps text on the researcher’s machine, which can make sensitive or restricted text data easier to use. The benchmark studies prompt-based classification only. For large labeled datasets with fixed label sets, fine-tuned local models, supervised classifiers, or embedding-based methods may be cheaper and more reliable. API endpoints can change behavior between releases even when the model name is stable, while local open-weight models are easier to freeze exactly.

References

- Ballard, Andrew O., Ryan DeTamble, Spencer Dorsey, Michael Heseltine, and Marcus Johnson. 2022. “Incivility in Congressional Tweets.” *American Politics Research* 50 (6): 769–80. <https://doi.org/10.1177/1532673X221109516>.
- Bestvater, Samuel E., and Burt L. Monroe. 2023. “Sentiment Is Not Stance: Target-Aware Opinion Classification for Political Text Analysis.” *Political Analysis* 31 (2): 235–56. <https://doi.org/10.1017/pan.2022.10>.
- Brandt, Patrick T., Sultan Alsarra, Vito D’Orazio, et al. 2025. “Extractive Versus Generative Language Models for Political Conflict Text Classification.” *Political Analysis*, 1–29. <https://doi.org/10.1017/pan.2025.10027>.
- Burnham, Michael. 2025. “Stance Detection: A Practical Guide to Classifying Political Beliefs in Text.” *Political Science Research and Methods* 13 (3): 611–28. <https://doi.org/10.1017/psrm.2024.35>.
- Burnham, Michael, Kayla Kahn, Ryan Yang Wang, and Rachel X. Peng. 2025. “Political DEBATE: Efficient Zero-Shot and Few-Shot Classifiers for Political Text.” *Political Analysis*, 1–15. <https://doi.org/10.1017/pan.2025.10028>.
- Chae, Youngjin, and Thomas Davidson. 2026. “Large Language Models for Text Classification: From Zero-Shot Learning to Instruction-Tuning.” *Sociological Methods & Research* 55 (2): 501–67. <https://doi.org/10.1177/00491241251325243>.
- Comparative Agendas Project. 2022. *Master Topics Codebook Version 1.1*. <https://www.comparativeagendas.net/pages/master-codebook>.
- Congressional Research Service. 2026. *CRS Reports*. <https://www.congress.gov/crs-products>.
- Di Cocco, Jessica, and Bernardo Monechi. 2022. “How Populist Are Parties? Measuring Degrees of Populism in Party Manifestos Using Supervised Machine Learning.” *Political Analysis* 30 (3): 311–27. <https://doi.org/10.1017/pan.2021.29>.
- Douglass, Rex W., Thomas Leo Scherer, J. Andrés Gannon, et al. 2024. “Introducing ICBc: An Event Extraction Dataset from Narratives about International Crises.” *Political Science Research and Methods* 12 (4): 729–49. <https://doi.org/10.1017/psrm.2024.17>.

- Erlich, Aaron, Stefano G. Dantas, Benjamin E. Bagozzi, Daniel Berliner, and Brian Palmer-Rubin. 2022. “Multi-Label Prediction for Political Text-as-Data.” *Political Analysis* 30 (4): 463–80. <https://doi.org/10.1017/pan.2021.15>.
- Garcia Corral, Paulina, Hanna Bechara, Ran Zhang, and Slava Jankin. 2024. “PolitiCause: An Annotation Scheme and Corpus for Causality in Political Texts.” *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024)* (Torino, Italia), 12836–45. <https://aclanthology.org/2024.lrec-main.1124/>.
- Gilardi, Fabrizio, Meysam Alizadeh, and Maël Kubli. 2023. “ChatGPT Outperforms Crowd Workers for Text-Annotation Tasks.” *Proceedings of the National Academy of Sciences* 120 (30): e2305016120. <https://doi.org/10.1073/pnas.2305016120>.
- González-Bustamante, Bastián. 2024. *Benchmarking LLMs in Political Content Text-Annotation: Proof-of-Concept with Toxicity and Incivility Data*. <https://arxiv.org/abs/2409.09741>.
- Halterman, Andrew, and Katherine A. Keith. 2026. “Codebook LLMs: Evaluating LLMs as Measurement Tools for Political Science Concepts.” *Political Analysis* 34 (2): 188–204. <https://doi.org/10.1017/pan.2025.10017>.
- Halterman, Andrew, Philip A. Schrodtt, Andreas Beger, Benjamin E. Bagozzi, and Grace I. Scarborough. 2023. *PLOVER and POLECAT: A New Political Event Ontology and Dataset*. <https://www.andrewhalterman.com/publication/plover-polecat-new-event-data/>.
- Haunss, Sebastian, Priska Daphi, Jan Matti Dollbaum, Lidiya Hristova, Pál Susánszky, and Elias Steinhilper. 2025. “PAPEA: A Modular Pipeline for the Automation of Protest Event Analysis.” *Political Science Research and Methods*, 1–18. <https://doi.org/10.1017/psrm.2025.10013>.
- Mellon, Jonathan, Jack Bailey, Ralph Scott, James Breckwoldt, Marta Miori, and Phillip Schmedeman. 2024. “Do AIs Know What the Most Important Issue Is? Using Language Models to Code Open-Text Social Survey Responses at Scale.” *Research & Politics* 11 (1): 1–7. <https://doi.org/10.1177/20531680241231468>.
- Müller, Stefan, and Naofumi Fujimura. 2025. “Campaign Communication and Legislative Leadership.” *Political Science Research and Methods* 13: 545–66. <https://doi.org/10.1017/psrm.2024.11>.
- Ornstein, Joseph T., Elise N. Blasingame, and Jake S. Truscott. 2025. “How to Train Your Stochastic Parrot: Large Language Models for Political Texts.” *Political Science Research and Methods* 13 (2): 264–81. <https://doi.org/10.1017/psrm.2024.64>.
- Osnabrügge, Moritz, Elliott Ash, and Massimo Morelli. 2023. “Cross-Domain Topic Classification for Political Texts.” *Political Analysis* 31 (1): 59–80. <https://doi.org/10.1017/pan.2021.37>.
- Pendzel, Sagi, Nir Lotan, Alon Zoizner, and Einat Minkov. 2023. *Detecting Multidimensional Political Incivility on Social Media*. <https://arxiv.org/abs/2305.14964>.
- Pipal, Christian, Eva-Maria Vogel, Morgan Wack, and Frank Esser. 2026. *Researchers Waste 80% of LLM Annotation Costs by Classifying One Text at a Time*. <https://arxiv.org/abs/2604.03684>.
- Rheault, Ludovic, Erica Rayment, and Andreea Musulan. 2019. “Politicians in the Line of Fire: Incivility and the Treatment of Women on Social Media.” *Research & Politics* 6 (1): 2053168018816228. <https://doi.org/10.1177/2053168018816228>.
- Sermpezis, Pavlos, Stelios Karamanidis, Eva Paraschou, et al. 2024. *AgoraSpeech: A Multi-Annotated Comprehensive Dataset of Political Discourse Through the Lens of Humans and AI*. <https://doi.org/10.5>

[281/zenodo.13957177](https://zenodo.org/doi/10.1177/2158244020919447).

Theocharis, Yannis, Pablo Barberá, Zoltán Fazekas, and Sebastian Adrian Popa. 2020. “The Dynamics of Political Incivility on Twitter.” *SAGE Open* 10 (2): 2158244020919447. <https://doi.org/10.1177/2158244020919447>.

Zhang, Meiqing, Furkan Cakmak, Markus Neumann, et al. 2025. “Comparable 2022 General Election Advertising Datasets from Meta and Google.” *Scientific Data* 12: 968. <https://doi.org/10.1038/s41597-025-05228-w>.

Companion materials

Prompt sources, batching details, summary CSVs, and reproduction instructions are on the GitHub repository at <https://github.com/hhilbig/polsci-open-bench>.

Tasks

Table 1: Thirty-four tasks in the benchmark. 'Type' is the report grouping used in Figure 3; it is a broad annotation type rather than a source-provided task family. 'Provenance': Verbatim = prompt unchanged from source paper, Verbatim-adapted = source content with format-only changes, Derived = our prompt reasoned from the published codebook. 'Prompt': short (15-35 lines) or long codebook (47-126 lines). This is the trait that determines local batching feasibility. The CMP task uses the Halterman & Keith (2026) domain-collapsed 7-class version of the CMP/MARPOR coding scheme, not the full CMP category set.

Task	Description	Source	Type	Labels	Provenance	Prompt
Gilardi relevance	Tweets, classified as about content moderation or not.	Gilardi 2023	Relevance & Harm	binary	Verbatim	short
Brandt relevance	News articles, classified as politics-relevant or not.	Brandt et al. 2025	Relevance & Harm	binary	Derived	short
Burnham PolNLI	Premise-hypothesis pairs, classified for entailment.	Burnham et al. 2025	Claims & Relations	binary	Derived	short
Ballard incivility	Tweets by U.S. members of Congress, classified for incivility.	Ballard 2022	Relevance & Harm	binary	Derived	short
Line of Fire incivility	Political tweets, classified for incivility.	Rheault et al. 2019	Relevance & Harm	binary	Derived	short
Theocharis incivility	Political tweets, classified for incivility.	Theocharis et al. 2020	Relevance & Harm	binary	Derived	short
Gilardi stance	Tweets about content moderation, classified by stance toward moderation.	Gilardi 2023	Position & Tone	3-class	Verbatim	short
SemEval stance	Tweets toward five named political targets, classified by stance.	Chae & Davidson 2026	Position & Tone	3-class	Derived	short
SCOTUS sentiment	Tweets about U.S. Supreme Court rulings, classified by sentiment.	Ornstein et al. 2025	Position & Tone	3-class	Derived	short
Wesleyan ad tone	U.S. Meta political ads, classified by tone toward a candidate.	Zhang et al. 2025	Position & Tone	3-class	Derived	short
CCC protest type	News stories about U.S. protest events, classified by event type.	Halterman & Keith 2026	Events & Actions	4-class	Verbatim-adapted	long codebook
BFRS violence	News stories about Pakistani political violence, classified by event type.	Halterman & Keith 2026	Events & Actions	12-class	Verbatim-adapted	long codebook
PAPEA protest forms	Protest snippets, classified by protest form.	Haunss et al. 2025	Events & Actions	7-class	Derived	long codebook
ICBe event type	Crisis-event sentences, classified by sentence event type.	Douglass et al. 2024	Events & Actions	5-class	Derived	short
CMP manifestos	Quasi-sentences from political party manifestos, classified by policy domain.	Halterman & Keith 2026	Issues & Topics	7-class	Derived	long codebook
Cross-domain topics	Political text, classified by broad policy domain.	Osnabruegge et al. 2023	Issues & Topics	8-class	Derived	long codebook
Campaign policy area	Campaign statements, classified by policy area.	Muller & Fujimura 2024	Issues & Topics	12-class	Derived	long codebook

Table 1: Thirty-four tasks in the benchmark. 'Type' is the report grouping used in Figure 3; it is a broad annotation type rather than a source-provided task family. 'Provenance': Verbatim = prompt unchanged from source paper, Verbatim-adapted = source content with format-only changes, Derived = our prompt reasoned from the published codebook. 'Prompt': short (15-35 lines) or long codebook (47-126 lines). This is the trait that determines local batching feasibility. The CMP task uses the Halterman & Keith (2026) domain-collapsed 7-class version of the CMP/MARPOR coding scheme, not the full CMP category set.

Task	Description	Source	Type	Labels	Provenance	Prompt
BES MII	Open-ended British Election Study responses, classified by issue topic.	Mellon et al. 2024	Issues & Topics	50-class	Verbatim-adapted	long codebook
Spanish protest toxicity	Spanish protest tweets, classified for toxic content.	toxicity-protests-ES	Relevance & Harm	binary	Derived	short
TwitCivility impoliteness	Political tweets, classified for impoliteness.	TwitCivility	Relevance & Harm	binary	Derived	short
GTD attack type	Global Terrorism Database events, classified by primary attack type.	Brandt et al. 2025	Events & Actions	9-class	Derived	long codebook
PAPEA protest claims	Protest snippets, classified by protest claim.	Haunss et al. 2025	Issues & Topics	28-class	Derived	long codebook
PLOVER CAMEO event	Gold-standard event records, classified by CAMEO event type.	PLOVER	Events & Actions	18-class	Derived	long codebook
PolNLI event entailment	Event-extraction PolNLI pairs, classified for entailment.	Burnham et al. 2025	Claims & Relations	binary	Derived	short
Women's March stance	Tweets about the Women's March, classified by stance.	Bestvater & Monroe	Position & Tone	binary	Derived	short
Trump stance	Political statements, classified by stance toward Trump.	Burnham et al. 2025	Position & Tone	3-class	Derived	short
COVID threat minimization	Political text, classified for COVID threat minimization.	Burnham et al. 2025	Claims & Relations	binary	Derived	short
Kavanaugh stance	Tweets about Brett Kavanaugh, classified by stance.	Bestvater & Monroe	Position & Tone	binary	Derived	short
ATI topics	Mexican access-to-information requests, classified by request topic.	Erlich et al.	Issues & Topics	7-label	Derived	long codebook
CAP party platforms	Party-platform quasi-statements, classified by CAP major topic.	CAP	Issues & Topics	21-class	Derived	long codebook
CAP CRS reports	Congressional Research Service titles and summaries, classified by CAP major topic.	CAP / CRS	Issues & Topics	21-class	Derived	long codebook
PolitiCAUSE causal relation	Political sentences, classified for whether they express a causal relation.	PolitiCAUSE	Claims & Relations	binary	Derived	short
Manifesto populism	Italian manifesto sentences, classified for populist rhetoric.	Di Cocco & Monechi	Claims & Relations	binary	Derived	short
AgoraSpeech criticism/agenda	Greek campaign-speech paragraphs, classified as criticism or agenda setting.	AgoraSpeech	Claims & Relations	2-class	Derived	short

Per-task winners

Table 2: Top model per task by the main F1 score across the unified 34-task benchmark, grouped by the report’s five annotation types. Accuracy and MCC refer to the same model-task pair. CIs and the full per-(task, model) table are on the GitHub repository.

Type	Task	Top model	Main F1	Accuracy	MCC
Relevance & Harm	Ballard incivility	Claude Sonnet 4.6	0.563	0.848	0.515
	Brandt relevance	Qwen3 14B	0.739	0.887	0.703
	Gilardi relevance	gpt-5.5	0.950	0.954	0.908
	Line of Fire incivility	DeepSeek V4 Pro	0.607	0.886	0.561
	Spanish protest toxicity	Gemma 4 26B	0.894	0.892	0.784
	Theocharis incivility	gpt-5.5	0.682	0.812	0.551
	TwitCivility impoliteness	Gemma 4 31B	0.631	0.794	0.503
Position & Tone	Gilardi stance	Claude Sonnet 4.6	0.568	0.648	0.471
	Kavanaugh stance	gpt-5.5	0.954	0.952	0.906
	SCOTUS sentiment	DeepSeek V4 Pro	0.668	0.718	0.543
	SemEval 2016 stance	gpt-5.5	0.773	0.768	0.659
	Trump stance	Claude Sonnet 4.6	0.804	0.826	0.730
	Wesleyan ad tone	gpt-5.4-nano	0.711	0.918	0.808
	Women’s March stance	Claude Sonnet 4.6	0.905	0.852	0.623
Events & Actions	BFRS violence (12 classes)	gpt-5.5	0.693	0.730	0.691
	CCC protest type	Qwen3 14B	0.499	0.610	0.387
	GTD attack type (9 classes)	gpt-5.4-nano	0.617	0.750	0.656
	ICBe event type (5 classes)	Gemma 4 26B	0.530	0.584	0.463
	PAPEA protest forms (7 classes)	gpt-5.5	0.965	0.972	0.958
	PLOVER CAMEO event (18 classes)	gpt-5.5	0.654	0.699	0.682
Claims & Relations	AgoraSpeech criticism/agenda	Mistral Sm 24B	0.922	0.924	0.843
	Burnham PolNLI	gpt-5.5	0.908	0.928	0.853
	COVID threat minimization	Mistral Sm 24B	0.602	0.819	0.499
	Manifesto populism	Mistral Sm 24B	0.455	0.904	0.403
	PolNLI event entailment	Gemma 4 31B	0.974	0.972	0.945
	PolitiCAUSE causal relation	Claude Sonnet 4.6	0.716	0.772	0.578
Issues & Topics	ATI topics (7 labels)	Claude Sonnet 4.6	0.502	NA	NA
	BES MII (50 classes)	gpt-5.5	0.547	0.946	0.934
	CAP CRS reports (21 classes)	gpt-5.5	0.732	0.774	0.757
	CAP party platforms (21 classes)	Claude Sonnet 4.6	0.731	0.736	0.719
	CMP manifestos (7 classes)	gpt-5.5	0.597	0.646	0.545
	Campaign policy area (12 classes)	gpt-5.5	0.636	0.750	0.691
	Cross-domain topics (8 classes)	Claude Sonnet 4.6	0.476	0.494	0.433
	PAPEA protest claims (28 classes)	gpt-5.5	0.462	0.636	0.606

Task-level uncertainty for model averages

Table 3: Paired task-level bootstrap intervals for cross-task mean-F1 differences among the four strongest models. Each bootstrap resamples the 34 tasks with replacement and recomputes the paired mean difference. These intervals assess sensitivity to benchmark composition; they do not replace the paired-by-item bootstrap intervals used for within-task comparisons.

Model comparison	Mean difference	95% task-bootstrap interval	Tasks
Claude Sonnet 4.6 - gpt-5.5	0.008	[-0.008, 0.024]	34
Claude Sonnet 4.6 - Gemma 4 31B	0.020	[0.005, 0.038]	34
Claude Sonnet 4.6 - DeepSeek V4 Pro	0.021	[0.004, 0.040]	34
gpt-5.5 - Gemma 4 31B	0.012	[-0.007, 0.033]	34
gpt-5.5 - DeepSeek V4 Pro	0.013	[-0.006, 0.032]	34
Gemma 4 31B - DeepSeek V4 Pro	0.001	[-0.020, 0.020]	34

Local runtime per 1,000 items

Figure 7 converts local median runtime per item into minutes per 1,000 items. The bars for 10 items per prompt use the full 34-task local batching grid, so they should be read alongside the response-format and label failure checks below.




Figure 7: Median local runtime per 1,000 items. One-at-a-time bars and bars for 10 items per prompt use the 34-task benchmark. Lower values are faster. Batching should be interpreted alongside response-format and label failure checks.

Model-selection upper bounds

Figure 2 in the main text compares the best local model with the best API model separately for each task. The largest API advantages appear on tasks that involve policy domains, event codes, causal relations, or protest claims. The largest local advantages appear on more focused tasks such as manifesto populism, COVID threat minimization, and CCC protest type.

Table 4: Mean F1 under single-model choices and task-specific upper bounds. The task-specific rows choose the best model after observing task-level performance, so they describe an upper bound rather than expected performance from a validation sample.

Selection rule	Model or rule	Mean F1
Best single model	Claude Sonnet 4.6	0.668
Best single local model	Gemma 4 31B	0.648
Best single API model	Claude Sonnet 4.6	0.668
Task-specific upper bound: local	Best local per task	0.673
Task-specific upper bound: API	Best API per task	0.688
Task-specific upper bound: all models	Best model per task	0.696

The best single model is Claude Sonnet 4.6, with a 34-task mean F1 of 0.668. The best single local model, Gemma 4 31B, reaches 0.648. If the local model can vary by task, the local task-specific upper bound rises to 0.673. The API upper bound reaches 0.688, and the all-model upper bound reaches 0.696. Because these rows choose after observing outcomes, they describe an upper bound on what validation can buy, not expected production performance.

Response-format or label failure rates above 5 percent

Table 5: Batched model-task pairs with response-format or label failure rates of at least 5 percent, not a complete list of all model-task pairs. Failures include responses that do not parse, responses with the wrong object or array shape, and responses that use labels outside the allowed label set. Local rows come from the full 34-task grid with 10 items per prompt for the five local models (170 pairs). OpenAI rows come from a separate 10-task prompt-batching diagnostic with 10 and 20 items per prompt (40 pairs); these rows do not enter the main one-item-at-a-time benchmark. Claude Sonnet and DeepSeek API prompt batching are not covered here.

Task	Model	Batch size	Scope	Failure rate
ATI topics (7 labels)	Gemma 4 26B	10	Local b=10	72%
CCC protest	gpt-5.5	10	OpenAI diagnostic	68%
CCC protest type	Qwen3 30B-A3B	10	Local b=10	51%
Wesleyan ads	gpt-5.5	20	OpenAI diagnostic	48%
CCC protest	gpt-5.5	20	OpenAI diagnostic	40%
BFRS Pakistan	gpt-5.5	20	OpenAI diagnostic	32%
Cross-domain topics (8 classes)	Qwen3 30B-A3B	10	Local b=10	30%
BFRS Pakistan	gpt-5.5	10	OpenAI diagnostic	28%
Wesleyan ads	gpt-5.5	10	OpenAI diagnostic	26%
BFRS violence (12 classes)	Qwen3 30B-A3B	10	Local b=10	18%
Ballard incivility	gpt-5.5	20	OpenAI diagnostic	16%
SemEval stance	gpt-5.5	20	OpenAI diagnostic	16%
Wesleyan ad tone	Qwen3 30B-A3B	10	Local b=10	16%
BFRS Pakistan	gpt-5.4-nano	10	OpenAI diagnostic	12%
SCOTUS sentiment	gpt-5.5	20	OpenAI diagnostic	12%
Ballard incivility	gpt-5.5	10	OpenAI diagnostic	10%
CMP manifestos	gpt-5.5	10	OpenAI diagnostic	10%
GTD attack type (9 classes)	Qwen3 30B-A3B	10	Local b=10	10%
CCC protest	gpt-5.4-nano	20	OpenAI diagnostic	8%
BFRS Pakistan	gpt-5.4-nano	20	OpenAI diagnostic	8%
PLOVER CAMEO event (18 classes)	Qwen3 30B-A3B	10	Local b=10	6%
CCC protest	gpt-5.4-nano	10	OpenAI diagnostic	6%

Performance under alternative metrics

Table 6: Unified 34-task model averages. Coverage is complete for all nine models in this version. The main F1 score is the task-specific summary metric; accuracy and MCC describe overall agreement. For imbalanced tasks, all three answer different questions.

Model	Coverage	Missing	Mean F1	Mean accuracy	Mean MCC
Claude Sonnet 4.6	34/34		0.668	0.777	0.632
gpt-5.5	34/34		0.660	0.775	0.627
Gemma 4 31B	34/34		0.648	0.770	0.615
DeepSeek V4 Pro	34/34		0.647	0.772	0.613
Gemma 4 26B	34/34		0.639	0.762	0.597
gpt-5.4-nano	34/34		0.632	0.748	0.591
Mistral Sm 24B	34/34		0.629	0.740	0.583
Qwen3 14B	34/34		0.608	0.749	0.570
Qwen3 30B-A3B	34/34		0.593	0.744	0.555